



Ciências  
ULisboa

# Concurrency Anomalies in PostgreSQL

ADVANCED DATABASES

**GROUP 1:** DANIELA VIEIRA, JOÃO RAIMUNDO, JOÃO RATO, MARIA VIEIRA


**PROFESSOR:** CÁTIA PESQUITA

22 November 2021

# From *PostgreSQL* to *PL/pgSQL*

- ‘Translate’ the queries
- Procedural Language that facilitate concurrency testing

## QUERY 1



Update to 1980/01/01 the release date of all albums for the genre Math rock, which were released in the 90's with an abstract over 200 characters, and that had most sales.

# SELECT FUNCTION – QUERY 1

```
CREATE OR REPLACE FUNCTION get_album_id_Q1()
    RETURNS SETOF INT AS $$

BEGIN
    RETURN QUERY
        SELECT albums.album_id
        FROM ((albums
            INNER JOIN bands ON albums.band_id = bands.band_id)
            INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
            INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
        WHERE genres.genre_name = 'Math rock'
        AND albums.release_date >= '1990/01/01'
        AND albums.release_date <= '1999/12/31'
        AND LENGTH(albums.abstract) > 200
        GROUP BY albums.album_id
        ORDER BY albums.sales
        DESC
        LIMIT 1;

END;
$$ LANGUAGE plpgsql;
```

# UPDATE FUNCTION – QUERY 1

```
CREATE OR REPLACE FUNCTION update_albums_release_date_Q1(  
    release_date_update_Q1 albums.release_date%TYPE)  
    RETURNS varchar AS $$  
DECLARE  
    album_id_Q1 albums.album_id%TYPE;  
BEGIN  
    SELECT get_album_id_Q1() INTO album_id_Q1;  
    UPDATE albums SET release_date = release_date_update_Q1 WHERE (albums.album_id = album_id_Q1);  
    RETURN 'UPDATED SUCCESSFULLY';  
END;  
$$ LANGUAGE plpgsql;
```

RUN TRANSACTION TWICE,  
IN DISTINCT SHELLS,  
AT THE SAME TIME

```
\set AUTOCOMMIT off
BEGIN;
SELECT update_albums_release_date_Q1('1980-01-01');
SELECT albums.album_id, albums.release_date
FROM (((albums
  INNER JOIN bands ON albums.band_id = bands.band_id)
  INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
  INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
WHERE genres.genre_name = 'Math rock'
AND albums.release_date = '1980-01-01'
AND LENGTH(albums.abstract) > 200
GROUP BY albums.album_id
ORDER BY albums.sales
DESC;
COMMIT;

SELECT pg_sleep(3);

SELECT albums.album_id
FROM (((albums
  INNER JOIN bands ON albums.band_id = bands.band_id)
  INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
  INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
WHERE genres.genre_name = 'Math rock'
AND albums.release_date >= '1990/01/01'
AND albums.release_date <= '1999/12/31'
AND LENGTH(albums.abstract) > 200
GROUP BY albums.album_id
ORDER BY albums.sales
DESC
LIMIT 5;
\set AUTOCOMMIT on
COMMIT;
```

# OUTPUTS

- EQUAL OUTPUTS
- UPDATING THE SAME PIECE OF DATA
- CONSISTENCY ANOMALY OCCURRED !
  - **DIRTY READ !**
- Manifests when a transaction can read uncommitted changes of some other concurrent transaction.

## SHELL 1 Transaction 1

```
update_albums_release_date_q1
-----
UPDATED SUCCESSFULLY
(1 row)
```

```
album_id | release_date
-----+-----
      23907 | 1980-01-01
(1 row)
```

```
album_id
-----
      21300
      2379
      2381
      34011
      2377
(5 rows)
```

## SHELL 2 Transaction 2

```
update_albums_release_date_q1
-----
UPDATED SUCCESSFULLY
(1 row)
```

```
album_id | release_date
-----+-----
      23907 | 1980-01-01
(1 row)
```

```
album_id
-----
      21300
      2379
      2381
      34011
      2377
(5 rows)
```

# SOLVING DIRTY READ PHENOMENA

Approaches that we used:

- Lock-Based Concurrency Control
- Isolation Levels



# SOLVING DIRTY READ PHENOMENA - LOCKS

- Implementing **LOCKS** into the Update Function -> **NEW FUNCTION** update\_albums\_release\_date\_Q1\_lock()
- **LOCK TABLE IN SHARE ROW EXCLUSIVE MODE**
- Protects a table **against concurrent data changes (reading & writing)**. Only one session can hold it at time

```
CREATE OR REPLACE FUNCTION update_albums_release_date_Q1_lock(  
    release_date_update_Q1 albums.release_date%TYPE)  
    RETURNS varchar AS $$  
DECLARE  
    album_id_Q1 albums.album_id%TYPE;  
BEGIN  
    LOCK TABLE albums IN SHARE ROW EXCLUSIVE MODE;  
    SELECT get_album_id_Q1() INTO album_id_Q1;  
    UPDATE albums SET release_date = release_date_update_Q1 WHERE(albums.album_id = album_id_Q1);  
    RETURN 'UPDATED SUCCESSFULLY';  
END;  
$$ LANGUAGE plpgsql;
```

# SOLVING DIRTY READ PHENOMENA - LOCKS - TRANSACTIONS

RUN TRANSACTIONS  
AGAIN TWICE,  
IN DISTINCT SHELLS  
AT THE SAME TIME,  
BUT THIS TIME WITH  
LOCKS IMPLEMENTED.

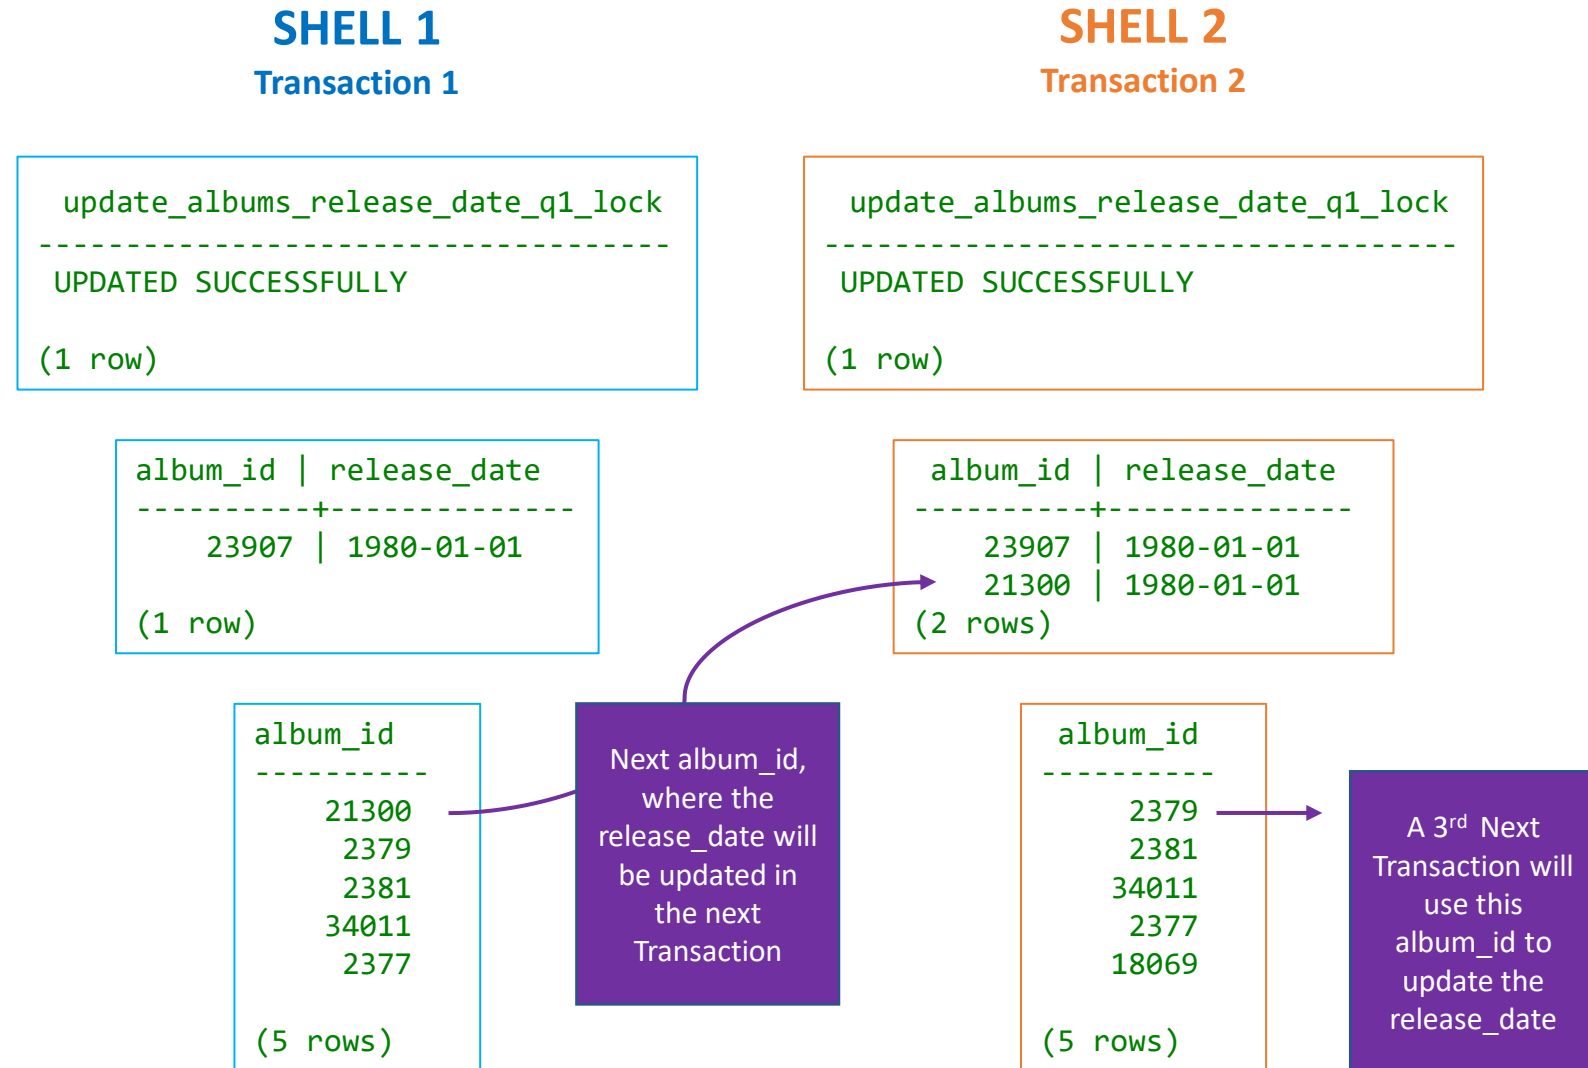
```
\set AUTOCOMMIT off
BEGIN;
SELECT update_albums_release_date_Q1_lock('1980-01-01');
SELECT pg_sleep(10);
SELECT albums.album_id, albums.release_date
FROM (((albums
  INNER JOIN bands ON albums.band_id = bands.band_id)
  INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
  INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
WHERE genres.genre_name = 'Math rock'
AND albums.release_date = '1980-01-01'
AND LENGTH(albums.abstract) > 200
GROUP BY albums.album_id
ORDER BY albums.sales
DESC;
COMMIT;

SELECT pg_sleep(5);

SELECT albums.album_id
FROM (((albums
  INNER JOIN bands ON albums.band_id = bands.band_id)
  INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
  INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
WHERE genres.genre_name = 'Math rock'
AND albums.release_date >= '1990/01/01'
AND albums.release_date <= '1999/12/31'
AND LENGTH(albums.abstract) > 200
GROUP BY albums.album_id
ORDER BY albums.sales
DESC
LIMIT 5;
\set AUTOCOMMIT on
COMMIT;
```

# SOLVING DIRTY READ PHENOMENA - LOCKS - OUTPUTS

- DIRTY READ PHENOMENA  
**SOLVED**
- The **1<sup>st</sup> Transaction**, occurred in **SHELL 1**, Updating the release\_date of the album\_id = '23907'
- The **2<sup>nd</sup> Transaction**, occurred in **SHELL 2**, Updating the release\_date on the album\_id = '21300'
- If a 3<sup>rd</sup> Transaction runs, it will Update the release\_date for the album\_id = '2379'. The **next** album with the **most sales**.



# SOLVING DIRTY READ PHENOMENA - ISOLATION LEVELS

## TRANSACTIONS

- TRANSACTION ISOLATION LEVEL USED: **REPEATABLE READ**
- Even though PG documentation refers that levels less restricted, like committed reads, does not allow Dirty Reads Phenomena – In our case it still happened, so we used a more restricted level.

```
\set AUTCOMMIT off
BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
SELECT update_albums_release_date_Q1('1980-01-01');
SELECT pg_sleep(10);
SELECT albums.album_id, albums.release_date
FROM (((albums
  INNER JOIN bands ON albums.band_id = bands.band_id)
  INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
  INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
WHERE genres.genre_name = 'Math rock'
AND albums.release_date = '1980-01-01'
AND LENGTH(albums.abstract) > 200
GROUP BY albums.album_id
ORDER BY albums.sales
DESC;

COMMIT;

SELECT pg_sleep(5);

SELECT albums.album_id
FROM (((albums
  INNER JOIN bands ON albums.band_id = bands.band_id)
  INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
  INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
WHERE genres.genre_name = 'Math rock'
AND albums.release_date >= '1990/01/01'
AND albums.release_date <= '1999/12/31'
AND LENGTH(albums.abstract) > 200
GROUP BY albums.album_id
ORDER BY albums.sales
DESC
LIMIT 5;

\set AUTCOMMIT on
COMMIT;
```

# SOLVING DIRTY READ PHENOMENA - ISOLATION LEVELS

## OUTPUTS

### SHELL 1 Transaction 1

```
update_albums_release_date_q1_lock
-----
UPDATED SUCCESSFULLY

(1 row)
```

album_id	release_date
23907	1980-01-01
21300	1980-01-01
2379	1980-01-01

(3 rows)

album_id
2381
34011
2377
18069
18067

(5 rows)

### SHELL 2 Transaction 2

```
ERROR:  could not serialize access due to concurrent update
CONTEXT:  SQL statement "UPDATE albums SET release_date = release_date_update_Q1 WHERE
(albums.album_id = album_id_Q1)"
PL/pgSQL function update_albums_release_date_q1(date) line 7 at SQL statement
bands_db=!# SELECT pg_sleep(10);
ERROR:  current transaction is aborted, commands ignored until end of transaction block
```

# INDUCING PHANTOM READ PHENOMENA

**Phantom Read** – Occurs when **two queries** ran successively, within the **same transaction**, displays **different** sets of **results** due to the **insertion** or **deletion** of one or more **rows** between read statements.

## NEXT STEPS:

- CREATE A NEW SELECT FUNCTION
- CREATE AN INSERT FUNCTION

# INDUCING PHANTOM READ PHENOMENA - SELECT FUNCTION

```
-- CREATE SELECT FUNCTION RETURNS A band_id

CREATE OR REPLACE FUNCTION get_band_id_Q1()
  RETURNS SETOF INT AS $$
BEGIN
  RETURN QUERY
    SELECT albums.band_id
    FROM (((albums
      INNER JOIN bands ON albums.band_id = bands.band_id)
      INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
      INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
    WHERE genres.genre_name = 'Math rock'
    AND albums.release_date >= '1990/01/01'
    AND albums.release_date <= '1999/12/31'
    AND LENGTH(albums.abstract) > 200
    GROUP BY albums.album_id
    ORDER BY albums.sales
    DESC
    LIMIT 1;
END;
$$ LANGUAGE plpgsql;
```

# INDUCING PHANTOM READ PHENOMENA - INSERT FUNCTION

```
-- CREATE INSERT FUNCTION
```

```
CREATE OR REPLACE FUNCTION insert_new_album(  
    album_id_T albums.album_id%TYPE,  
    album_name_T albums.album_name%TYPE,  
    sales_T albums.sales%TYPE,  
    time_T albums.running_time%TYPE,  
    date_T albums.release_date%TYPE,  
    abstract_T albums.abstract%TYPE)  
    RETURNS varchar AS $$  
DECLARE  
    band_id_Q1 albums.band_id%TYPE;  
BEGIN  
    SELECT get_band_id_Q1() INTO band_id_Q1;  
    INSERT INTO albums (album_id,band_id,album_name,sales,running_time,release_date,abstract)  
    VALUES (album_id_T,band_id_Q1,album_name_T,sales_T,time_T,date_T,abstract_T);  
    RETURN 'INSERTED SUCCESSFULLY';  
END;  
$$ LANGUAGE plpgsql;
```



# INDUCING PHANTOM READ PHENOMENA - TRANSACTIONS

## SHELL 1 Transaction 1

```
\set AUTCOMMIT off
BEGIN;
SELECT album_id, band_id, release_date, sales
  FROM albums as A
 WHERE A.release_date >= '1998/08/10'
 AND A.release_date <= '1998/08/17'
 GROUP BY A.album_id, A.band_id, A.release_date, A.sales
 ORDER BY A.sales
 DESC;
SELECT pg_sleep(20);
SELECT album_id, band_id, release_date, sales
  FROM albums as A
 WHERE A.release_date >= '1998/08/10'
 AND A.release_date <= '1998/08/17'
 GROUP BY A.album_id, A.band_id, A.release_date, A.sales
 ORDER BY A.sales
 DESC;
\set AUTCOMMIT on
COMMIT;
```

## SHELL 2 Transaction 2

```
\set AUTCOMMIT off
BEGIN;
SELECT insert_new_album('34716', 'TEST ALBUM', '59', '30.0', '1998/08/16', 'TEST ALBUM ABSTRACT');
\set AUTCOMMIT on
COMMIT;
```

# INDUCING PHANTOM READ PHENOMENA - OUTPUTS

## 1<sup>st</sup> SELECT

## 2<sup>nd</sup> SELECT

### SHELL 1 Transaction 1

album_id	band_id	release_date	sales
11958	3074	1998-08-12	9951
8551	2149	1998-08-12	9880
9799	2540	1998-08-10	9636
4994	1277	1998-08-12	9368
21629	5510	1998-08-17	9176
11957	3074	1998-08-12	8466
15571	4083	1998-08-12	8205
4008	1088	1998-08-17	7657

(8 rows)

album_id	band_id	release_date	sales
11958	3074	1998-08-12	9951
8551	2149	1998-08-12	9880
9799	2540	1998-08-10	9636
4994	1277	1998-08-12	9368
21629	5510	1998-08-17	9176
11957	3074	1998-08-12	8466
15571	4083	1998-08-12	8205
4008	1088	1998-08-17	7657
34716	4705	1998-08-16	59

(9 rows)

### SHELL 2 Transaction 2

```
insert_new_album
-----
INSERTED SUCCESSFULLY
```

Transaction 2 influence the result of Transaction 1

-> **Phantom READ Phenomena**

# SOLVING PHANTOM READ PHENOMENA


Approaches that we used:

- Lock-Based Concurrency Control
- Isolation Levels

# SOLVING PHANTOM READ PHENOMENA - LOCKS

- Implementing **LOCKS** in **Transaction 1 (SHELL 1)**
- **LOCK TABLE** albums **IN EXCLUSIVE MODE**
- Prevents other transactions to **write in table albums** while Transaction 1 runs (until being COMMITTED)

```
\set AUTCOMMIT off
BEGIN;
LOCK TABLE albums IN EXCLUSIVE MODE;
SELECT album_id, band_id, release_date, sales
  FROM albums as A
 WHERE A.release_date >= '1998/08/10'
 AND A.release_date <= '1998/08/17'
 GROUP BY A.album_id, A.band_id, A.release_date, A.sales
 ORDER BY A.sales
 DESC;
SELECT pg_sleep(20);
SELECT album_id, band_id, release_date, sales
  FROM albums as A
 WHERE A.release_date >= '1998/08/10'
 AND A.release_date <= '1998/08/17'
 GROUP BY A.album_id, A.band_id, A.release_date, A.sales
 ORDER BY A.sales
 DESC;
\set AUTCOMMIT on
COMMIT;
```



# SOLVING PHANTOM READ PHENOMENA - LOCKS - OUTPUTS

## 1<sup>st</sup> SELECT

## 2<sup>nd</sup> SELECT

### SHELL 1 Transaction 1

album_id	band_id	release_date	sales
11958	3074	1998-08-12	9951
8551	2149	1998-08-12	9880
9799	2540	1998-08-10	9636
4994	1277	1998-08-12	9368
21629	5510	1998-08-17	9176
11957	3074	1998-08-12	8466
15571	4083	1998-08-12	8205
4008	1088	1998-08-17	7657

(8 rows)

album_id	band_id	release_date	sales
11958	3074	1998-08-12	9951
8551	2149	1998-08-12	9880
9799	2540	1998-08-10	9636
4994	1277	1998-08-12	9368
21629	5510	1998-08-17	9176
11957	3074	1998-08-12	8466
15571	4083	1998-08-12	8205
4008	1088	1998-08-17	7657


(8 rows)

- **Transaction 1** - Same set of results in both queries
- **Transaction 2** - Waits until Transaction 1 being committed to start the transaction.

# SOLVING PHANTOM READ PHENOMENA - ISOLATION LEVELS

- Implemented in **Transaction 1**
- TRANSACTION ISOLATION LEVEL USED: **REPEATABLE READ**

```
\set AUTOCOMMIT off
BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ;
SELECT album_id, band_id, release_date, sales
    FROM albums as A
    WHERE A.release_date >= '1998/08/10'
    AND A.release_date <= '1998/08/17'
    GROUP BY A.album_id, A.band_id, A.release_date, A.sales
    ORDER BY A.sales
    DESC;
SELECT pg_sleep(20);
SELECT album_id, band_id, release_date, sales
    FROM albums as A
    WHERE A.release_date >= '1998/08/10'
    AND A.release_date <= '1998/08/17'
    GROUP BY A.album_id, A.band_id, A.release_date, A.sales
    ORDER BY A.sales
    DESC;
\set AUTOCOMMIT on
COMMIT;
```



# SOLVING PHANTOM READ PHENOMENA - ISOLATION LEVELS

## OUTPUTS

### SHELL 1 Transaction 1

#### 1<sup>st</sup> SELECT

album_id	band_id	release_date	sales
11958	3074	1998-08-12	9951
8551	2149	1998-08-12	9880
9799	2540	1998-08-10	9636
4994	1277	1998-08-12	9368
21629	5510	1998-08-17	9176
11957	3074	1998-08-12	8466
15571	4083	1998-08-12	8205
4008	1088	1998-08-17	7657

(8 rows)

#### 2<sup>nd</sup> SELECT

album_id	band_id	release_date	sales
11958	3074	1998-08-12	9951
8551	2149	1998-08-12	9880
9799	2540	1998-08-10	9636
4994	1277	1998-08-12	9368
21629	5510	1998-08-17	9176
11957	3074	1998-08-12	8466
15571	4083	1998-08-12	8205
4008	1088	1998-08-17	7657


(8 rows)

### SHELL 2 Transaction 2

```
insert_new_album
-----
INSERTED SUCCESSFULLY
```

- **Transaction 1** - Same set of results in both queries
- **Transaction 2** - It can be committed during Transaction 1 process, without affecting its set of results. Did not report any errors.

## QUERY 2



Update to 0 the sales from the album with the most sales in the first decade of the year 2000, and which the running time is longer than 45 minutes.



# SELECT FUNCTION – QUERY 2

```
CREATE OR REPLACE FUNCTION get_album_id_most_sales_Q2()  
| RETURNS SETOF INT AS $$  
BEGIN  
| RETURN QUERY  
|     SELECT album_id  
|     FROM albums  
|     WHERE running_time > '45'  
|     AND release_date >= '2000/01/01'  
|     AND release_date <= '2010/12/31'  
|     ORDER BY sales  
|     DESC  
|     LIMIT 1;  
END;  
$$ LANGUAGE plpgsql;
```

## UPDATE FUNCTION – QUERY 2

```
CREATE OR REPLACE FUNCTION update_sales_Q2()  
| RETURNS varchar AS $$  
DECLARE  
| album_id_before albums.album_id%TYPE;  
BEGIN  
|     SELECT get_album_id_most_sales() INTO album_id_before;  
|     PERFORM pg_sleep(10);  
|     UPDATE albums SET sales = 0 WHERE (albums.album_id = album_id_before);  
|     RETURN 'UPDATE SUCCESSFULL';  
END;  
$$ LANGUAGE plpgsql;
```

# RUN TRANSACTION, IN DISTINCT SHELLS, AT THE SAME TIME

```
\set AUTOCOMMIT off
BEGIN;
    SELECT update_sales_Q2();
    SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 0;
COMMIT;

SELECT pg_sleep(3);

SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 0;
\set AUTOCOMMIT on
COMMIT;
```

# OUTPUTS

- EQUAL OUTPUTS
- UPDATING THE SAME PIECE OF DATA
- CONSISTENCY ANOMALY OCCURRED !
  - **DIRTY READ !**

Update function  
output

## SHELL 1 Transaction 1

```
update_sales_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

First SELECT query  
output

band_id	sales	release_date	running_time
464	0	2015-06-11	76.933334
595	0	2002-12-11	42.416668
2700	0	1991-06-14	51.216667

(3 row)

Second SELECT query  
output

band_id	sales	release_date	running_time
464	0	2015-06-11	76.933334
595	0	2002-12-11	42.416668
2700	0	1991-06-14	51.216667

(3 row)

## SHELL 2 Transaction 2

```
update_sales_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

band_id	sales	release_date	running_time
464	0	2015-06-11	76.933334
595	0	2002-12-11	42.416668
2700	0	1991-06-14	51.216667

(3 row)

band_id	sales	release_date	running_time
464	0	2015-06-11	76.933334
595	0	2002-12-11	42.416668
2700	0	1991-06-14	51.216667

(3 row)

# SOLVING DIRTY READ PHENOMENA - LOCKS

```
CREATE OR REPLACE FUNCTION update_sales_lock_Q2()  
| RETURNS varchar AS $$  
DECLARE  
| album_id_before albums.album_id%TYPE;  
BEGIN  
|     LOCK TABLE albums;  
|     SELECT get_album_id_most_sales() INTO album_id_before;  
|     PERFORM pg_sleep(10);  
|     UPDATE albums SET sales = 0 WHERE (albums.album_id = album_id_before);  
|     RETURN 'Update SUCCESSFULL';  
END;  
$$ LANGUAGE plpgsql;
```

# SOLVING DIRTY READ PHENOMENA - LOCKS - TRANSACTIONS

RUN TRANSACTIONS IN DISTINCT SHELLS AT THE SAME TIME,  
BUT THIS TIME WITH **LOCKS IMPLEMENTED**.

```
\set AUTOCOMMIT off
BEGIN;
    SELECT update_sales_lock_Q2();
    SELECT band_id,album_id, sales, release_date, running_time FROM albums WHERE sales = 0;
COMMIT;

SELECT pg_sleep(3);

SELECT band_id,album_id, sales, release_date, running_time FROM albums WHERE sales = 0;
\set AUTOCOMMIT on
COMMIT;
```

# SOLVING DIRTY READ PHENOMENA - LOCKS - OUTPUTS

## SHELL 1 Transaction 1

Update function  
output

```
update_sales_lock_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

First SELECT query  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  |    76.933334
(1 row)
```

Second SELECT query  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  |    76.933334
(1 row)
```

## SHELL 2 Transaction 2

```
update_sales_lock_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  |    76.933334
(1 row)
```

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  |    76.933334
    464 |    1338 |     0 | 2015-06-11  |    79.133333
(2 row)
```

# SOLVING DIRTY READ PHENOMENA - ISOLATION LEVELS

TRANSACTION ISOLATION LEVEL USED: **REPEATABLE READ**

```
\set AUTOCOMMIT off
BEGIN TRANSACTION ISOLATION LEVEL Read Committed;
    SELECT update_sales_Q2();
    SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 0;
COMMIT;

SELECT pg_sleep(10);

SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 0;
\set AUTOCOMMIT on
COMMIT;
```



# SOLVING DIRTY READ PHENOMENA - ISOLATION LEVELS

## OUTPUTS

### SHELL 1 Transaction 1

Update function  
output

```
update_sales_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

First SELECT query  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  | 76.933334
(1 row)
```

Second SELECT query  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  | 76.933334
(1 row)
```

### SHELL 2 Transaction 2

```
update_sales_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  | 76.933334
(1 row)
```

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  | 76.933334
    464 |    1338 |     0 | 2015-06-11  | 79.13333
(2 row)
```

# INDUCING LOST UPDATE PHENOMENA

**Lost update** – occur when the second update overwrites the first one.

## UPDATE FUNCTION

```
CREATE OR REPLACE FUNCTION update_sales_lost_update_Q2()  
| RETURNS varchar AS $$  
DECLARE  
| album_id_before albums.album_id%TYPE;  
BEGIN  
|     SELECT get_album_id_most_sales() INTO album_id_before;  
|     PERFORM pg_sleep(10);  
|     UPDATE albums SET sales = 1 WHERE (albums.album_id = album_id_before);  
|     RETURN 'UPDATE SUCCESSFUL';  
END;  
$$ LANGUAGE plpgsql;
```

# INDUCING LOST UPDATE PHENOMENA - OUTPUTS

## SHELL 1 Transaction 1

Update function  
output

```
update_sales_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

First SELECT query  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  |    76.933334
(1 row)
```

Second SELECT query  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  |    76.933334
(1 row)
```

## SHELL 2 Transaction 2

```
update_sales_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  |    76.933334
(1 row)
```

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11  |    76.933334
    464 |    1338 |     0 | 2015-06-11  |    79.13333
(2 row)
```

# SOLVING LOST UPDATE PHENOMENA - LOCKS

## SHELL 1

### Transaction 1

```
\set AUTOCOMMIT off
BEGIN;
    SELECT update_sales_lock_Q2();
    SELECT band_id,album_id, sales, release_date, running_time FROM albums WHERE sales = 0;
COMMIT;

SELECT pg_sleep(10);

SELECT band_id,album_id, sales, release_date, running_time FROM albums WHERE sales = 0;
\set AUTOCOMMIT on
```

## SHELL 2

### Transaction 2

```
\set AUTOCOMMIT off
BEGIN;
    SELECT update_sales_lost_update_Q2();
    SELECT band_id, sales, release_date, running_time FROM albums WHERE sales= 0 or sales = 1;
COMMIT;

SELECT pg_sleep(10);

SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 0 or sales = 1;
\set AUTOCOMMIT on
```

# SOLVING LOST UPDATE PHENOMENA - LOCKS - OUTPUTS

## SHELL 1 Transaction 1

Update function  
output

```
update_sales_lock_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

First SELECT query  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
      464 |      1280 |      0 | 2015-06-11   |      76.933334
(1 row)
```

Second SELECT query  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
      464 |      1280 |      0 | 2015-06-11   |      76.933334
(1 row)
```

## SHELL 2 Transaction 2

```
update_sales_lost_update_Q2
-----
UPDATED SUCCESSFULLY
(1 row)
```

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
      464 |      1280 |      0 | 2015-06-11   |      76.933334
(1 row)
```

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
      464 |      1280 |      0 | 2015-06-11   |      76.933334
      464 |      1338 |      1 | 2015-06-11   |      79.13333
(2 row)
```

# SOLVING LOST UPDATE PHENOMENA - ISOLATION LEVELS

## SHELL 1 Transaction 1

```
\set AUTOCOMMIT off
BEGIN TRANSACTION ISOLATION LEVEL Repeatable read;
    SELECT update_sales_Q2();
    SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 0;
COMMIT;

SELECT pg_sleep(10);

SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 0;
\set AUTOCOMMIT on
COMMIT;
```

## SHELL 2 Transaction 2

```
\set AUTOCOMMIT off
BEGIN;
    SELECT update_sales_lost_update_Q2();
    SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 1;
COMMIT;

SELECT pg_sleep(10);

SELECT band_id, sales, release_date, running_time FROM albums WHERE sales = 1;
\set AUTOCOMMIT on
COMMIT;
```

# SOLVING LOST UPDATE PHENOMENA - ISOLATION LEVELS

## OUTPUTS

### SHELL 1 Transaction 1

```
update_sales_Q2
-----
UPDATED SUCCESSFULLY

(1 row)
```

Update function  
output

```
band_id | album_id | sales | release_date | running_time
-----+-----+-----+-----+-----
    464 |    1280 |     0 | 2015-06-11   |      76.933334

(1 row)
```

First SELECT query  
output

```
album_id
-----
    2381
    34011
    2377
    18069
    18067
(5 rows)
```

Second SELECT query  
output

### SHELL 2 Transaction 2

```
ERROR:  could not serialize access due to concurrent update
CONTEXT:  SQL statement "UPDATE albums SET release_date = release_date_update_Q1 WHERE
(albums.album_id = album_id_Q1)"
PL/pgSQL function update_albums_release_date_q1(date) line 7 at SQL statement
bands_db=!# SELECT pg_sleep(10);
ERROR:  current transaction is aborted, commands ignored until end of transaction block
```

# FINAL REMARKS

- *Lock-Based Concurrency Control or Isolation Levels* implementations can sometimes have a negative impact on the database performance
- The most coherent approach to use in our case is the explicit locks, as it does not demonstrate any implementation concerns in contrast to the isolation levels, and proved to be the best option to ensure a better availability of the database
- It is necessary to evaluate the trade-offs between data consistency and concurrency, considering what use the database will be given.

