



Ciências
ULisboa

Indexing and Optimization

ADVANCED DATABASES


GROUP 1: DANIELA VIEIRA, JOÃO RAIMUNDO, JOÃO RATO, MARIA VIEIRA

PROFESSOR: CÁTIA PESQUITA

13 dezembro 2021

PostgreSQL - Query Optimization

QUERY 1



Update to 1980/01/01 the release date of all albums for the genre Math rock, which were released in the 90's with an abstract over 200 characters, and that had most sales.

Query 1 – Execution Plan Without Indexes

```
EXPLAIN ANALYZE SELECT albums.band_id, albums.release_date, albums.sales, genres.genre_name, albums.abstract
FROM (((albums
INNER JOIN bands ON albums.band_id = bands.band_id)
INNER JOIN bands_genre ON bands.band_id = bands_genre.band_id)
INNER JOIN genres ON bands_genre.genre_id = genres.genre_id)
    WHERE genres.genre_name = 'Math rock'
    AND albums.release_date >= '1990/01/01'
    AND albums.release_date <= '1999/12/31'
    AND LENGTH(albums.abstract) > 200
GROUP BY albums.band_id, albums.release_date, albums.abstract, albums.sales, genres.genre_name
ORDER BY albums.sales
DESC
LIMIT 1;
```

Query 1 – Execution Plan - Output

QUERY PLAN

```
-----
Limit (cost=3563.78..3563.80 rows=1 width=504) (actual time=148.644..148.673 rows=1 loops=1)
  -> Group (cost=3563.78..3563.93 rows=10 width=504) (actual time=148.638..148.664 rows=1 loops=1)
    Group Key: albums.sales, albums.band_id, albums.release_date, albums.abstract, genres.genre_name
    -> Sort (cost=3563.78..3563.81 rows=10 width=504) (actual time=148.631..148.655 rows=1 loops=1)
      Sort Key: albums.sales DESC, albums.band_id, albums.release_date, albums.abstract
      Sort Method: quicksort  Memory: 35kB
      -> Nested Loop (cost=415.34..3563.62 rows=10 width=504) (actual time=83.918..148.549 rows=14 loops=1)
        Join Filter: (albums.band_id = bands.band_id)
        -> Hash Join (cost=415.06..3560.10 rows=11 width=508) (actual time=83.871..148.236 rows=14 loops=1)
          Hash Cond: (albums.band_id = bands_genre.band_id)
          -> Seq Scan on albums (cost=0.00..3136.76 rows=2179 width=492) (actual time=8.642..81.148 rows=4841 loops=1)
            Filter: ((release_date >= '1990-01-01'::date) AND (release_date <= '1999-12-31'::date) AND (length(abstract) > 200))
            Rows Removed by Filter: 29247
          -> Hash (cost=414.51..414.51 rows=44 width=16) (actual time=61.220..61.232 rows=62 loops=1)
            Buckets: 1024  Batches: 1  Memory Usage: 11kB
            -> Hash Join (cost=10.66..414.51 rows=44 width=16) (actual time=0.270..61.105 rows=62 loops=1)
              Hash Cond: (bands_genre.genre_id = genres.genre_id)
              -> Seq Scan on bands_genre (cost=0.00..341.32 rows=23632 width=8) (actual time=0.060..33.324 rows=23632 loops=1)
              -> Hash (cost=10.65..10.65 rows=1 width=16) (actual time=0.142..0.146 rows=1 loops=1)
                Buckets: 1024  Batches: 1  Memory Usage: 9kB
                -> Seq Scan on genres (cost=0.00..10.65 rows=1 width=16) (actual time=0.021..0.111 rows=1 loops=1)
                  Filter: ((genre_name)::text = 'Math rock'::text)
                  Rows Removed by Filter: 531
            -> Index Only Scan using bands_pkey on bands (cost=0.29..0.31 rows=1 width=4) (actual time=0.017..0.017 rows=1 loops=14)
              Index Cond: (band_id = bands_genre.band_id)
              Heap Fetches: 0
        -> Index Only Scan using bands_pkey on bands (cost=0.29..0.31 rows=1 width=4) (actual time=0.017..0.017 rows=1 loops=14)
          Index Cond: (band_id = bands_genre.band_id)
          Heap Fetches: 0
```

Planning Time: 86.578 ms
Execution Time: 148.778 ms
(28 rows)



Planning Time: 86.578 ms
Execution Time: 148.778 ms
(28 rows)

Query 1 – Query Plan Output with Index

```
CREATE INDEX date_abstract ON albums USING btree (release_date, length(abstract));
```


```
--- Output Query 1 Plan with data_abstract index
```

QUERY PLAN

```
-----
Limit (cost=3112.42..3112.43 rows=1 width=504) (actual time=83.892..83.926 rows=1 loops=1)
  -> Group (cost=3112.42..3112.57 rows=10 width=504) (actual time=83.885..83.916 rows=1 loops=1)
        Group Key: albums.sales, albums.band_id, albums.release_date, albums.abstract, genres.genre_name
  -> Sort (cost=3112.42..3112.44 rows=10 width=504) (actual time=83.879..83.908 rows=1 loops=1)
        Sort Key: albums.sales DESC, albums.band_id, albums.release_date, albums.abstract
        Sort Method: quicksort Memory: 35kB
  -> Nested Loop (cost=569.88..3112.25 rows=10 width=504) (actual time=63.279..83.706 rows=14 loops=1)
        Join Filter: (albums.band_id = bands.band_id)
        -> Hash Join (cost=569.59..3108.73 rows=11 width=508) (actual time=63.207..83.384 rows=14 loops=1)
              Hash Cond: (albums.band_id = bands_genre.band_id)
              -> Bitmap Heap Scan on albums (cost=154.53..2685.40 rows=2179 width=492) (actual time=1.638..17.599 rows=4841 loops=1)
                    Recheck Cond: ((release_date >= '1990-01-01'::date) AND (release_date <= '1999-12-31'::date) AND (length(abstract) > 200))
                    Heap Blocks: exact=474
              -> Bitmap Index Scan on date_abstract (cost=0.00..153.99 rows=2179 width=0) (actual time=1.496..1.499 rows=4841 loops=1)
                    Index Cond: ((release_date >= '1990-01-01'::date) AND (release_date <= '1999-12-31'::date) AND (length(abstract) > 200))
        -> Hash (cost=414.51..414.51 rows=44 width=16) (actual time=59.980..59.992 rows=62 loops=1)
              Buckets: 1024 Batches: 1 Memory Usage: 11kB
              -> Hash Join (cost=10.66..414.51 rows=44 width=16) (actual time=0.340..59.831 rows=62 loops=1)
                    Hash Cond: (bands_genre.genre_id = genres.genre_id)
                    -> Seq Scan on bands_genre (cost=0.00..341.32 rows=23632 width=8) (actual time=0.076..30.948 rows=23632 loops=1)
                    -> Hash (cost=10.65..10.65 rows=1 width=16) (actual time=0.177..0.180 rows=1 loops=1)
                          Buckets: 1024 Batches: 1 Memory Usage: 9kB
                          -> Seq Scan on genres (cost=0.00..10.65 rows=1 width=16) (actual time=0.029..0.156 rows=1 loops=1)
                                Filter: ((genre_name)::text = 'Math rock'::text)
                                Rows Removed by Filter: 531
        -> Index Only Scan using bands_pkey on bands (cost=0.29..0.31 rows=1 width=4) (actual time=0.018..0.018 rows=1 loops=14)
              Index Cond: (band_id = bands_genre.band_id)
              Heap Fetches: 0
```

```
Planning Time: 30.279 ms
Execution Time: 84.312 ms
(30 rows)
```

**B-tree Composite
Index**



Planning Time: 30.279 ms
Execution Time: 84.312 ms
(30 rows)

Query 1 – Clustered Index


```
CREATE INDEX date ON albums USING btree (release_date);
```

```
CLUSTER albums USING date;  
ANALYZE albums;
```

```
--- EXPLAIN ANALYZE QUERY 1 WITH the albums table clustered by release date
```


QUERY PLAN

```
-----  
Limit (cost=1126.69..1126.71 rows=1 width=506) (actual time=80.687..80.716 rows=1 loops=1)  
  -> Group (cost=1126.69..1127.02 rows=22 width=506) (actual time=80.682..80.708 rows=1 loops=1)  
        Group Key: albums.sales, albums.band_id, albums.release_date, albums.abstract, genres.genre_name  
        -> Sort (cost=1126.69..1126.75 rows=22 width=506) (actual time=80.677..80.701 rows=1 loops=1)  
              Sort Key: albums.sales DESC, albums.band_id, albums.release_date, albums.abstract  
              Sort Method: quicksort  Memory: 35kB  
        -> Nested Loop (cost=415.63..1126.20 rows=22 width=506) (actual time=58.888..80.652 rows=14 loops=1)  
              Join Filter: (albums.band_id = bands.band_id)  
              -> Hash Join (cost=415.35..1118.20 rows=25 width=510) (actual time=58.871..80.525 rows=14 loops=1)  
                    Hash Cond: (albums.band_id = bands_genre.band_id)  
                    -> Index Scan using date on albums (cost=0.29..684.44 rows=4924 width=494) (actual time=0.047..17.814 rows=4841 loops=1)  
                          Index Cond: ((release_date >= '1990-01-01'::date) AND (release_date <= '1999-12-31'::date))  
                          Filter: (length(abstract) > 200)  
                          Rows Removed by Filter: 1717  
                    -> Hash (cost=414.51..414.51 rows=44 width=16) (actual time=56.719..56.731 rows=62 loops=1)  
                          Buckets: 1024  Batches: 1  Memory Usage: 11kB  
                          -> Hash Join (cost=10.66..414.51 rows=44 width=16) (actual time=0.218..56.609 rows=62 loops=1)  
                                Hash Cond: (bands_genre.genre_id = genres.genre_id)  
                                -> Seq Scan on bands_genre (cost=0.00..341.32 rows=23632 width=8) (actual time=0.074..28.020 rows=23632 loops=1)  
                                -> Hash (cost=10.65..10.65 rows=1 width=16) (actual time=0.086..0.090 rows=1 loops=1)  
                                      Buckets: 1024  Batches: 1  Memory Usage: 9kB  
                                      -> Seq Scan on genres (cost=0.00..10.65 rows=1 width=16) (actual time=0.029..0.075 rows=1 loops=1)  
                                            Filter: ((genre_name)::text = 'Math rock'::text)  
                                            Rows Removed by Filter: 531  
              -> Index Only Scan using bands_pkey on bands (cost=0.29..0.31 rows=1 width=4) (actual time=0.003..0.004 rows=1 loops=14)  
                    Index Cond: (band_id = bands_genre.band_id)  
                    Heap Fetches: 0  
Planning Time: 1.158 ms  
Execution Time: 80.790 ms  
(29 rows)
```



Planning Time: 1.158 ms
Execution Time: 80.790 ms

QUERY 2



Update to 0 the sales from the album with the most sales in the first decade of the year 2000, and which the running time is longer than 45 minutes.

Query 2 – Execution Plan Without Indexes

```
EXPLAIN ANALYSE SELECT A.band_id, A.sales, A.release_date, A.running_time
FROM albums AS A
  WHERE A.running_time >= '45'
  AND A.release_date >= '2000/01/01'
  AND A.release_date <= '2010/12/31'
GROUP BY  A.band_id,A.sales, A.release_date, A.running_time
ORDER BY A.sales
DESC
LIMIT 1;
```

Query 2 – Execution Plan - Output

QUERY PLAN

```
-----
Limit (cost=3231.48..3231.48 rows=1 width=16) (actual time=105.643..105.653 rows=1 loops=1)
  -> Sort (cost=3231.48..3247.17 rows=6276 width=16) (actual time=105.638..105.644 rows=1 loops=1)
        Sort Key: sales DESC
        Sort Method: top-N heapsort  Memory: 25kB
        -> HashAggregate (cost=3137.34..3200.10 rows=6276 width=16) (actual time=82.943..94.541 rows=9291 loops=1)
              Group Key: sales, band_id, release_date, running_time
              Batches: 1  Memory Usage: 913kB
              -> Seq Scan on albums a (cost=0.00..3051.54 rows=8580 width=16) (actual time=20.549..68.347 rows=9291 loops=1)
                    Filter: ((running_time >= '45'::real) AND (release_date >= '2000-01-01'::date) AND (release_date <= '2010-12-31'::date))
                    Rows Removed by Filter: 24797
Planning Time: 1.245 ms
Execution Time: 105.984 ms
(12 rows)
```



Planning Time: 1.245 ms
Execution Time: 105.984 ms
(12 rows)

Query 2 – Query Plan Output with Index

**B-tree Composite
Index**

```
CREATE INDEX sales_time_date ON albums USING btree (sales,running_time,release_date);
```

QUERY PLAN

```
-----  
Limit (cost=2.03..3.82 rows=1 width=16) (actual time=0.848..0.859 rows=1 loops=1)  
  -> Group (cost=2.03..11258.45 rows=6276 width=16) (actual time=0.842..0.850 rows=1 loops=1)  
        Group Key: sales, band_id, release_date, running_time  
        -> Incremental Sort (cost=2.03..11172.65 rows=8580 width=16) (actual time=0.839..0.844 rows=1 loops=1)  
              Sort Key: sales DESC, band_id, release_date, running_time  
              Presorted Key: sales  
              Full-sort Groups: 1  Sort Method: quicksort  Average Memory: 26kB  Peak Memory: 26kB  
              -> Index Scan Backward using sales_time_date on albums a (cost=0.29..10862.49 rows=8580 width=16) (actual time=0.087..0.717 rows=36 loops=1)  
                    Index Cond: ((running_time >= '45'::real) AND (release_date >= '2000-01-01'::date) AND (release_date <= '2010-12-31'::date))  
Planning Time: 13.946 ms  
Execution Time: 1.019 ms  
(11 rows)
```



Planning Time: 13.946 ms
Execution Time: 1.019 ms
(11 rows)

Query 2 – Clustered Index

QUERY PLAN

```
Limit (cost=2.03..3.83 rows=1 width=16) (actual time=0.918..0.929 rows=1 loops=1)
  -> Group (cost=2.03..11257.76 rows=6270 width=16) (actual time=0.912..0.920 rows=1 loops=1)
    Group Key: sales, band_id, release_date, running_time
    -> Incremental Sort (cost=2.03..11172.02 rows=8574 width=16) (actual time=0.909..0.914 rows=1 loops=1)
      Sort Key: sales DESC, band_id, release_date, running_time
      Presorted Key: sales
      Full-sort Groups: 1 Sort Method: quicksort Average Memory: 26kB Peak Memory: 26kB
      -> Index Scan Backward using sales_time_date on albums a (cost=0.29..10862.08 rows=8574 width=16) (actual time=0.090..0.776 rows=36 loops=1)
        Index Cond: ((running_time >= '45'::real) AND (release_date >= '2000-01-01'::date) AND (release_date <= '2010-12-31'::date))
Planning Time: 7.104 ms
Execution Time: 1.089 ms
(11 rows)
```



Planning Time: 7.104 ms
Execution Time: 1.089 ms

NoSQL - Query Optimization

- Permission problems between the MongoDB Shell (mongosh) and MongoDB Compass
- We used the MongoDB Compass GUI tabs to create the indexes and the explain query plans

Query 1 – Execution Plan Without Index – Output

Query Performance Summary

Documents Returned: 1
Index Keys Examined: 0
Documents Examined: 34088
Actual Query Execution Time (ms): 27
Sorted in Memory: yes
No index available for this query.

```
{
  "stage": "SORT",
  "nReturned": 1,
  "executionTimeMillisEstimate": 3,
  "works": 34092,
  "advanced": 1,
  "needTime": 34090,
  "needYield": 0,
  "saveState": 34,
  "restoreState": 34,
  "isEOF": 1,
  "sortPattern": {
    "sales": -1
  },
  "memLimit": 33554432,
  "limitAmount": 1,
  "type": "simple",
  "totalDataSizeSorted": 24526,
  "usedDisk": false
}
```

```
{
  "stage": "COLLSCAN",
  "filter": {
    "$and": [
      {
        "genres": {
          "$eq": "Math rock"
        }
      },
      {
        "release_date": {
          "$lte": "1999-12-31T00:00:00.000Z"
        }
      },
      {
        "abstract": {
          "$regex": "^[\\s\\S]{200,}$"
        }
      }
    ]
  },
  "nReturned": 22,
  "executionTimeMillisEstimate": 3,
  "works": 34090,
  "advanced": 22,
  "needTime": 34067,
  "needYield": 0,
  "saveState": 34,
  "restoreState": 34,
  "isEOF": 1,
  "direction": "forward",
  "docsExamined": 34088
}
```

Query 1 – Execution Plan With Index - Output

```
db.albums.createIndex({"sales":-1}, {"release_date":1}, name: "sales_date")
```

Query Performance Summary:

```
Documents Returned: 1
Index Keys Examined: 1926
Documents Examined: 1926
Actual Query Execution Time (ms): 5
Sorted in Memory: no
Query used the following index: sales_date
```

```
{
  "stage": "LIMIT",
  "nReturned": 1,
  "executionTimeMillisEstimate": 4,
  "works": 1927,
  "advanced": 1,
  "needTime": 1925,
  "needYield": 0,
  "saveState": 1,
  "restoreState": 1,
  "isEOF": 1,
  "limitAmount": 1
}
```

```
{
  "stage": "FETCH",
  "filter": {
    "$and": [
      {
        "genres": {
          "$eq": "Math rock"
        }
      },
      {
        "release_date": {
          "$lte": "1999-12-31T00:00:00.000Z"
        }
      }
    ]
  },
  "abstract": {
    "$regex": "^[[\\s\\S]]{200,}$"
  }
},
  "nReturned": 1,
  "executionTimeMillisEstimate": 4,
  "works": 1926,
  "advanced": 1,
  "needTime": 1925,
  "needYield": 0,
  "saveState": 1,
  "restoreState": 1,
  "isEOF": 0,
  "docsExamined": 1926,
  "alreadyHasObj": 0
}
```

```
{
  "stage": "IXSCAN",
  "nReturned": 1926,
  "executionTimeMillisEstimate": 1,
  "works": 1926,
  "advanced": 1926,
  "needTime": 0,
  "needYield": 0,
  "saveState": 1,
  "restoreState": 1,
  "isEOF": 0,
  "keyPattern": {
    "sales": -1,
    "release_date": 1
  },
  "indexName": "sales_date",
  "isMultiKey": false,
  "multiKeyPaths": {
    "sales": [],
    "release_date": []
  },
  "isUnique": false,
  "isSparse": false,
  "isPartial": false,
  "indexVersion": 2,
  "direction": "forward",
  "indexBounds": {
    "sales": [
      "[MaxKey, MinKey]"
    ],
    "release_date": [
      "[MinKey, MaxKey]"
    ]
  },
  "keysExamined": 1926,
  "seeks": 1,
  "dupsTested": 0,
  "dupsDropped": 0
}
```


Query 2 – Execution Plan Without Index – Output

Query Performance Summary:

Documents Returned: 1
Index Keys Examined: 0
Documents Examined: 34088
Actual Query Execution Time (ms): 28
Sorted in Memory: yes
No index available for this query.

```
{
  "stage": "SORT",
  "nReturned": 1,
  "executionTimeMillisEstimate": 9,
  "works": 34092,
  "advanced": 1,
  "needTime": 34090,
  "needYield": 0,
  "saveState": 34,
  "restoreState": 34,
  "isEOF": 1,
  "sortPattern": {
    "sales": -1
  },
  "memLimit": 33554432,
  "limitAmount": 1,
  "type": "simple",
  "totalDataSizeSorted": 18924046,
  "usedDisk": false
}
```

```
{
  "stage": "COLLSCAN",
  "filter": {
    "$and": [
      {
        "release_date": {
          "$lte": "2010-12-31T00:00:00.000Z"
        }
      }
    ],
    {
      "running_time": {
        "$gte": 45
      }
    }
  ],
  "nReturned": 14628,
  "executionTimeMillisEstimate": 8,
  "works": 34090,
  "advanced": 14628,
  "needTime": 19461,
  "needYield": 0,
  "saveState": 34,
  "restoreState": 34,
  "isEOF": 1,
  "direction": "forward",
  "docsExamined": 34088
}
```

Query 2 – Execution Plan With Index - Output

```
db.albums.createIndex({"sales":-1}, {"release_date":1}, {"running_time":1}, name: "sales_date_time")
```

```
Query Performance Summary
Documents Returned: 1
Index Keys Examined: 1
Documents Examined: 1
Actual Query Execution Time (ms): 1
Sorted in Memory: no
Query used the following index: sales_date_time
```

```
{
  "stage": "LIMIT",
  "nReturned": 1,
  "executionTimeMillisEstimate": 1,
  "works": 2,
  "advanced": 1,
  "needTime": 0,
  "needYield": 0,
  "saveState": 0,
  "restoreState": 0,
  "isEOF": 1,
  "limitAmount": 1
}
```

```
{
  "stage": "FETCH",
  "filter": {
    "$and": [
      {
        "release_date": {
          "$lte": "2010-12-31T00:00:00.000Z"
        }
      },
      {
        "running_time": {
          "$gte": 45
        }
      }
    ]
  },
  "nReturned": 1,
  "executionTimeMillisEstimate": 1,
  "works": 1,
  "advanced": 1,
  "needTime": 0,
  "needYield": 0,
  "saveState": 0,
  "restoreState": 0,
  "isEOF": 0,
  "docsExamined": 1,
  "alreadyHasObj": 0
}
```

```
{
  "stage": "IXSCAN",
  "nReturned": 1,
  "executionTimeMillisEstimate": 1,
  "works": 1,
  "advanced": 1,
  "needTime": 0,
  "needYield": 0,
  "saveState": 0,
  "restoreState": 0,
  "isEOF": 0,
  "keyPattern": {
    "sales": -1,
    "release_date": 1,
    "running_time": 1
  },
  "indexName": "sales_date_time",
  "isMultiKey": false,
  "multiKeyPaths": {
    "sales": [],
    "release_date": [],
    "running_time": []
  },
  "isUnique": false,
  "isSparse": false,
  "isPartial": false,
  "indexVersion": 2,
  "direction": "forward",
  "indexBounds": {
    "sales": [
      "[MaxKey, MinKey]"
    ],
    "release_date": [
      "[MinKey, MaxKey]"
    ],
    "running_time": [
      "[MinKey, MaxKey]"
    ]
  },
  "keysExamined": 1,
  "seeks": 1,
  "dupsTested": 0,
  "dupsDropped": 0
}
```

Discussion

- NoSQL relatively faster than PostgreSQL (query execution times)
 - Query 1 needs 3 joins in PostgreSQL – Takes time!
 - NoSQL was structured in a way that it does not need multiple aggregations, only one -> Albums
- No alterations for RDB, **but** we know a few:
 - Denormalization (genre names) – reducing the number of joins
 - Horizontal decomposition (transforming *release_date* into *decades* and creating new relations [90s, 80s, etc])
 - Not performed since the queries execution times were fast enough (ms)
- No alterations needed for NoSQL

